An attractor account of numerosity perception

Journal Title
XX(X):1–7
©The Author(s) 2026
Reprints and permission:
sagepub.co.uk/journalsPermissions.nav
DOI: 10.1177/ToBeAssigned
www.sagepub.com/
SAGE

Eudald Correig-Fraga[1]

Abstract

Numerical cognition is classically split between two systems: an exact mechanism for small sets and an approximate one for larger ones, separated by a poorly understood boundary. We propose a mechanistic alternative: the integers are attractor memories on the number line, each a well that widens with magnitude; numerosity judgement is pattern completion of a noisy estimate. Fitted to individual brief-exposure responses, the model outperforms non-attractor and Weber-style alternatives; its dynamics improve the fit without adding parameters and predict response times from convergence time and residual uncertainty. The model predicts its own limit: the numerosity where neighbouring wells stop being resolvable, which grows with exposure and marks the response-time plateau. A single core number system, the full set of integer attractors, is read in two modes: exact where the wells resolve (the Exact Number System) and approximate where they merge (the Approximate Number System), a single mechanism across both regimes.

Keywords

numerical cognition, subitizing, exact number system, attractor network, approximate number system

Introduction

Humans perceive small numerosities (1–3/4) quickly and accurately, a phenomenon known as subitizing (Kaufman et al. 1949). Beyond this range, response time increases and accuracy deteriorates, with variability that grows with set size in the ratio-dependent manner characteristic of Weber's law (Mandler and Shebo 1982; Dehaene 2011). These two signatures have been documented across ages, cultures, and species (Feigenson et al. 2004; Pica et al. 2004; Nieder 2016), and they anchor the standard view of numerical cognition as the product of two core systems: an Object Tracking System (OTS) that individuates a handful of items exactly, and an Approximate Number System (ANS) that represents larger magnitudes noisily (Piazza et al. 2011; Feigenson et al. 2002; Halberda et al. 2008).

What the two-systems picture does not provide is a mechanistic account of the boundary between them. The limit of exact perception is treated as a property of the OTS (i.e., a fixed capacity of three or four items) yet empirically it is not fixed: the range of near-exact responding expands with viewing time (Cheyette and Piantadosi 2020) and contracts in developmental dyscalculia (Decarli et al. 2020). A boundary that moves with exposure and experience is difficult to reconcile with a hand-off between two independent modules, and no current account derives it from the structure of the underlying representation: it simply marks where one explanation stops and another begins.

Computational work has shown that the two regimes need not reflect two mechanisms. Cheyette and Piantadosi (2020) offer the clearest normative formulation: under a power-law prior and a finite information-processing constraint, near-exact performance for small numerosities and Weber-like variability for larger ones emerge from a single optimization. Cheyette et al. (2024) later pushed the argument upstream, proposing that the psychophysics of number follows from bandwidth limits on visuospatial memory rather than from number-specific representations, and Malo et al. (2023) showed that a small quantum spin model can encode numerosity in the spectrum of its collective dynamics. Recurrent networks with biologically grounded on-centre off-surround dynamics likewise produce a partition of the numerosity range from a single architecture (Sengupta et al. 2014), and image-computable networks account for how numerosity is extracted from the visual scene in the first place (Stoianov and Zorzi 2012; Testolin et al. 2020). These accounts establish that unification is possible at the computational level; what is still missing is a mechanistic model of the representation itself: one that says what makes small numbers exact, derives rather than assumes where exactness ends, and yields the approximate regime from the same structure.

Here we propose such a model. The mental number line is cast as a one-dimensional attractor landscape in the sense of modern Hopfield networks (Hopfield 1982; Ramsauer et al. 2021): each integer is a stored memory, a well whose width grows super-linearly with magnitude, parameterized by the frequency with which numbers occur in the environment (Dehaene and Mehler 1992); numerosity judgement is pattern completion, a noisy sensory estimate relaxing to the nearest resolvable well. We evaluate the

[1]Innovamat Education, Avinguda de la Generalitat, 217, Catalonia (Spain)

Corresponding author:
Eudald Correig-Fraga, Innovamat Education, Avinguda de la Generalitat, 217, Catalonia (Spain).
Email: eudald.correig@innovamat.com

model on ∼53,000 individual responses from the brief-exposure dataset of Cheyette and Piantadosi (2020), where it outperforms Weber/ANS and normative alternatives; its clean-up dynamics improve the fit without adding parameters and, together with the residual readout uncertainty, generate response times. The same landscape also predicts its own limit: the numerosity n_{crit} at which neighbouring wells stop being resolvable, which grows with exposure duration and coincides with the onset of the empirical response-time plateau. On this view, a single representation, the full lattice of integer attractors that we take to be the core number system, is read in two modes. When the wells are resolvable, pattern completion recovers the exact numerosity, the Exact Number System (ENS); beyond n_{crit} the wells have merged and the same readout returns a coarse, ratio-dependent estimate, the Approximate Number System (ANS). Both are estimates of the one landscape, so a single mechanism spans the exact and approximate regimes of numerosity estimation rather than handing off between independent modules. The classical subitizing (Object Tracking System) range corresponds to the fully resolved lower part of this system, so the proposal refines rather than replaces the core-systems picture (Feigenson et al. 2004). As an exploratory extension, we apply the model to exact-enumeration data from children with developmental dyscalculia (Decarli et al. 2020), where it localizes the deficit as a shorter exact range (Supplementary Material 5).

Methods

The model maps a noisy perceptual estimate of numerosity onto a discrete response through relaxation on an energy landscape whose stable states sit at the integers (Fig. 1a). A single trial has three stages: a sensory estimate provides a starting point, gradient descent cleans it up toward a well, and a probability distribution over integers is read out from the resting state. The same dynamics also determine how long the trial takes.

Attractor landscape. Each integer n is a stored memory represented by a Gaussian well centred at n, with a width that grows with magnitude. We tie the width to how often each number occurs in the environment: number-word frequency follows an approximately power-law (Benford-like) decline across languages (Dehaene and Mehler 1992), so frequent small numbers occupy narrow, stable basins and rare large numbers occupy broad ones. Writing $f_n = \log_{10}(1 + 1/n)$ for the frequency of n, the well width is

$$\sigma_n = (\sigma_c / f_n)^{\beta}, \tag{1}$$

where σ_c sets the overall scale and the exponent β controls how strongly frequency modulates width. We leave β free because there is no a priori reason for the mapping from frequency to width to be linear; as we show below, the data require it to be super-linear ($\beta \approx 2$), which is what makes small numbers sharply resolved and large ones broadly overlapping. Each well contributes a log-likelihood

$$g_n(x) = -\log \sigma_n - \frac{(x-n)^2}{2\,\sigma_n^2}, \tag{2}$$

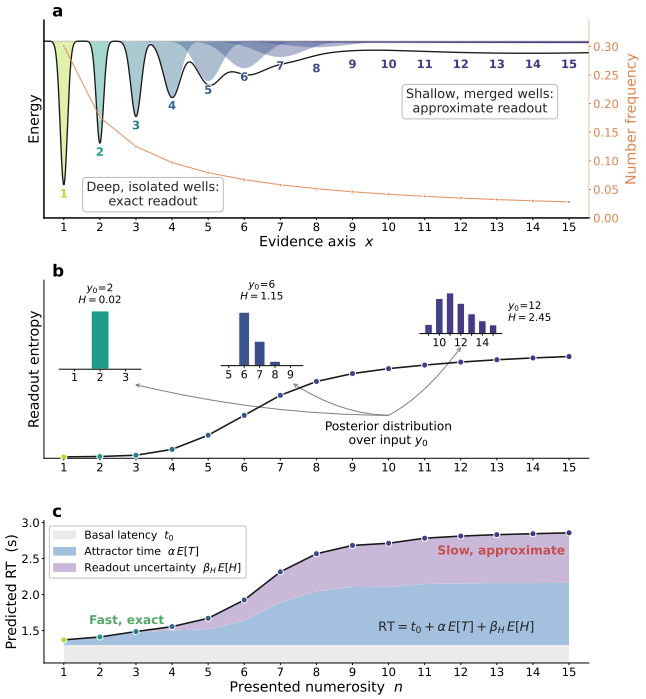


Figure 1. The attractor model of the number line (parameters fitted at 160 ms). a, Energy landscape over the internal evidence axis x: each integer is a well whose width grows with magnitude (well widths set by number frequency f_n, orange curve), so small numbers sit in deep, isolated basins that clean up to an exact value, whereas beyond n_{crit} (dotted line, where adjacent wells reach $\sigma_n = 1$) the wells merge into a shallow, approximate code. **b**, Expected readout entropy $E[H]$ per presented numerosity, before (grey) and after (black) clean-up, with the attractor's expected convergence time $E[T]$ (blue, right axis); insets show example post-clean-up posteriors. Clean-up removes the residual uncertainty of the small-n readout at the cost of a few settling steps; near the boundary the settling time grows and the residual entropy rises. **c**, The resulting response time, $\text{RT} = t_0 + \tau\, E[T] + \kappa\, E[H]$, decomposed into its basal, convergence-time, and readout-uncertainty components.

and, following modern Hopfield networks (Ramsauer et al. 2021), the attraction potential is the log-sum-exp over all wells, $V(x) = -\log \sum_n \exp\big(g_n(x)\big)$.

Clean-up and readout. A sensory estimate y_0 initialises the state; we model the sensory channel as $y_0 \sim \mathcal{N}\big(n,\ s_{00}^2 + s_0\, n^{\lambda_0}\big)$, a neutral front end whose noise grows with numerosity (s_{00} fixed, s_0 and λ_0 fitted). From $x = y_0$ the state descends the potential, $x_{t+1} = x_t - \eta\, \partial V/\partial x$. Because V is a log-sum-exp over the stored integers, each step moves the state toward a softmax combination of the wells, weighted by their precisions, so every stored integer contributes to the final position. The descent runs until the state settles or begins to oscillate, giving a resting state x^* reached in T steps. A distribution over integers is then read out, $p(n \mid x^*) = \text{softmax}\big(g_n(x^*)\big)$, from which we take the mean estimate and the most probable integer (exact accuracy).

Response time. The chronometric prediction combines the two quantities the mechanism itself produces: the time the state takes to settle and the uncertainty that remains once

it has,

$$\mathrm{RT} = t_0 + \tau\,\mathbb{E}[T] + \kappa\,\mathbb{E}\big[H(p(n \mid x^*))\big], \qquad (3)$$

where T is the number of clean-up steps, $H = -\sum_n p \log p$ is the residual readout entropy, and expectations are taken over the sensory noise. The convergence-time term $\tau\,\mathbb{E}[T]$ captures the deliberation the attractor dynamics require (long near the boundary, where the landscape flattens and the state settles slowly), while $\kappa\,\mathbb{E}[H]$ captures how undecided the readout remains between neighbouring integers. We compare this hybrid link against its two nested special cases (entropy only, $\tau = 0$; steps only, $\kappa = 0$). The basal latency t_0 is task-specific and fixed per dataset.

Trial-level likelihood. We fit the model to the individual responses of Cheyette and Piantadosi (2020), in whom adult participants ($N = 100$ per condition) briefly viewed dot arrays of 1–15 items and typed an estimate, at five exposure durations (40, 80, 160, 320, 640 ms). Rather than fitting summary curves, we evaluate the categorical likelihood of the ∼53,000 individual integer reports. For each numerosity the response probability marginalises the readout over the sensory noise,

$$p(r \mid n) = \int p\big(r \mid x^*(y_0)\big)\,\phi\big(y_0; n, \sigma(n)\big)\,dy_0, \qquad (4)$$

computed by deterministic grid quadrature (Gauss–Hermite fails on the near-step clean-up integrand; see Supplementary Material 1). Typed responses run well beyond the well support, so we add a uniform lapse channel, $p_{\mathrm{final}}(r \mid n) = (1-\ell)\,p(r \mid n) + \ell\,/\,R_{\max}$, with ℓ fitted per duration and counted in all information criteria. The response model therefore has four mechanism parameters (σ_c, s_0, λ_0, β) plus the lapse, fitted independently per duration by differential evolution (Storn and Price 1997) and Nelder–Mead polishing; the RT channel adds τ and κ. The response, RT, and coupled joint likelihoods are optimised as described in Supplementary Material 1.

Model comparison. To isolate the contribution of the attractor dynamics, we compare the model against a graded set of alternatives fitted under the identical trial-level convention: static (no-clean-up) readouts of the same wells, a Weber/ANS code with magnitude-proportional noise, and a log-prior normative proxy in the spirit of Cheyette and Piantadosi (2020). Because turning clean-up on or off leaves the parameter count unchanged, the attractor–vs–static contrast is a zero-parameter comparison. Supplementary Material 2 gives the full ladder and the role of the width exponent. Model quality is summarised by AIC and BIC over the full response count table.

Resolution limit. Finally, the landscape yields a parameter-free prediction of where the exact code breaks down. Adjacent wells stop being reliably distinguishable when their separation falls to one standard deviation, $d' = \mathrm{spacing}/\sigma_n = 1$ (Macmillan and Creelman 2004); solving $\sigma_n = 1$ gives the critical numerosity n_{crit} at which the discrete code loses resolution. This quantity is derived from the fitted landscape, not from any boundary in the data.

Results

The clean-up dynamics improve the fit at equal parameters. At the trial level the models separate sharply (Fig. 2a; full ladder in Supplementary Material 2). Switching the clean-up dynamics off while holding everything else fixed gives a matched, zero-parameter control: the static readout of the same wells fits worse by $\Delta\mathrm{AIC} \approx 1260$, so the dynamics sharpen the response distribution without adding flexibility. The model also outperforms the standard alternatives by a wide margin under both AIC and BIC: the Weber/ANS code by ≈ 4400 and the normative log-prior proxy by ≈ 7700 (Supplementary Material 2). The width exponent is fitted freely and lands firmly super-linear ($\beta \approx 2$), so the wells are sharp for small numbers and broadly overlapping for large ones, a property the data determine rather than one we impose (Supplementary Material 2). The clean-up advantage is concentrated at the briefest exposures, where the sensory evidence is poorest; at 320–640 ms a static readout of the same landscape does about as well on the response distribution (Fig. 2a), as expected if the attractor has largely settled during the exposure itself, leaving little for the post-stimulus dynamics to resolve.

One landscape fits estimation, accuracy, and response time. With these ingredients the model reproduces all three behavioural signatures jointly (Fig. 3): pooled across durations, $R^2 = 0.99$ for the mean estimate, 0.94 for exact accuracy, and 0.88 for response time, the last a genuine cross-prediction of a channel fitted to the individual trial RTs rather than to a summary curve. The one systematic miss is a mild under-prediction of exact accuracy in the subitizing band ($n = 2$–4): the model captures the trend toward near-perfect small-number performance but not its full sharpness, which we return to in the Discussion.

Response time follows from the attractor dynamics. Because clean-up is a dynamical process, response time can be derived from how long it runs rather than from a static uncertainty proxy, and the dynamical link is the better one (Fig. 2b): a link based on the attractor's settling time alone already outperforms the entropy proxy by $\Delta\mathrm{AIC} \approx 540$, and the hybrid link that adds the residual entropy is best (≈770 better than entropy, ≈230 better than steps alone), a ranking the coupled response–RT fit preserves. The two terms play complementary roles (Fig. 1b–c): near the boundary the landscape flattens and the state settles slowly, so the convergence-time term drives most of the rise in RT, while the residual-entropy term dominates once the wells have merged. The chronometric prediction is thus not an add-on but a second readout of the same mechanism.

The model predicts its own resolution limit. Because the well width grows with numerosity, the landscape specifies where the discrete code must fail: the numerosity n_{crit} at which adjacent wells reach $\sigma_n = 1$ (Fig. 2c). This crossing falls *inside* the fitted range, at $n_{\mathrm{crit}} \approx 4.6, 7.5, 8.7, 9.2, 10.0$ for the five durations; it is not imposed by any cutoff, since the model is fitted on the full range. Two independent facts converge on it. First, n_{crit} grows monotonically with exposure: more viewing time sharpens the wells, so the exact range extends from ∼5 at a 40 ms flash to ∼10 at 640 ms. This duration dependence is a property of the landscape, not of a

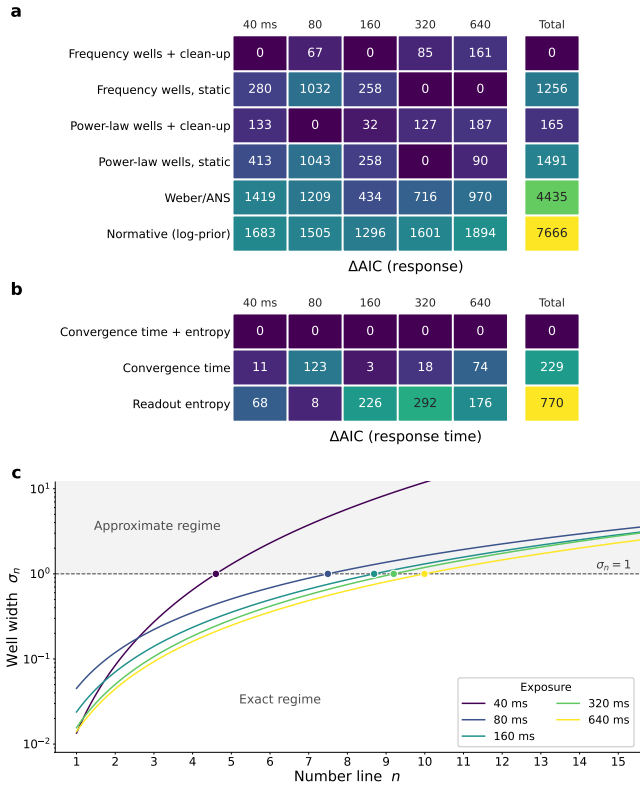


Figure 2. Trial-level model comparison, response time, and the predicted resolution limit. a, Response-likelihood ΔAIC over the $\sim$53,000 individual responses, per exposure duration and summed (Total), each cell relative to the best of its column (darker is better; shading scaled within each panel). Each well family (frequency-based width, Eq. (1), free β; or power-law width $\sigma_n \propto n^p$) appears with its clean-up dynamics and as a matched static readout of the same wells (no extra parameters): clean-up wins at the briefest exposures, the static readouts catch up at 320–640 ms, and in total the winning frequency-well attractor leads its static twin by ≈ 1260, with the Weber/ANS code and the normative log-prior proxy far behind. **b**, Response-time-channel ΔAIC on the winning landscape, same layout. The link combining the attractor's convergence time and the residual readout entropy is best at every exposure; convergence time alone stays close, while the static entropy proxy falls behind at 160–640 ms. **c**, The fitted width law σ_n per exposure duration. Wells are resolvable ($\sigma_n < 1$) at small n and merge ($\sigma_n \geq 1$) beyond the critical numerosity n_{crit} (dots), which grows from $\approx$4.6 at 40 ms to $\approx$10 at 640 ms: more exposure sharpens the wells and extends the exact range.

fixed capacity: the wells sharpen as the sensory evidence improves, exactly what a trainable memory should show (a constrained duration-independent landscape is examined in Supplementary Material 2). Second, the response-time plateau in the data begins near n_{crit}: within the exact range the model tracks RT closely, and above it the deliberation-time prediction keeps rising while the empirical RT flattens (Fig. 1c and Supplementary Material 4), the point at which participants stop resolving and start approximating. Accuracy and chronometry, two independent channels, thus mark the same boundary.

A parameter-recovery analysis (Supplementary Material 3) confirms that the quantities we interpret are identifiable at the trial level: the width exponent β and, especially, the derived n_{crit} recover almost perfectly ($R^2 = 0.89$ and 0.996), while the two sensory parameters trade off and are interpreted only through the combined noise profile $\sigma(n)$. A brief exploratory application to developmental dyscalculia, where the same model localises the deficit as a shorter exact range, is reported in Supplementary Material 5.

Discussion

We have described numerosity perception as pattern completion on an attractor landscape: the integers are memories whose widths grow with magnitude, and a noisy estimate relaxes to the nearest resolvable well. Fitted to individual responses, the model reproduces estimation, accuracy, and response time together, and, unlike accounts that stipulate a subitizing capacity, it derives the limit of exact perception from the structure of the representation itself. The exact and approximate regimes become two readings of one landscape rather than two systems joined at an unexplained seam.

Two features of the account carry the explanatory weight. The first is the clean-up dynamics. Turning clean-up on is a zero-parameter change, yet it improves the fit substantially, and the attractor's convergence time predicts response time better than a static uncertainty proxy. The clean-up is associative retrieval in the modern Hopfield sense: each step moves the state toward a softmax-weighted average of the stored integers, so every well shapes the outcome. Because those patterns are the ordered points of the number line, not arbitrary memories, the softmax competition among neighbours is sharp for small, well-separated numbers and broad for large, overlapping ones. The graded loss of resolution and the lengthening of settling time toward the boundary both follow from this structure. The second feature is that the well widths are set by number frequency. This ties the representation to the broader view of numbers as memory traces whose strength reflects experience (Anderson and Schooler 1991; Verguts and Fias 2004): frequent small numbers are sharply encoded, rare large ones diffusely.

Casting the code as a frequency-shaped memory has a developmental consequence that a fixed-capacity account lacks. The resolution limit n_{crit} is not a constant but a property of how sharp the wells are, and we already see it move with exposure over milliseconds; over development it should move with numerical experience. As children encounter number words and counting, the wells that carry the most exposure should sharpen, extending the exact range and pushing the transition to approximate perception upward. This predicts that the boundary sits at very small numbers early in development and grows with numeracy, and, more specifically, that it should be trained by exposure to *numbers*, not by general improvement in visual acuity: an intervention that sharpens the number code should differ in its signature from one that improves perception at large (Butterworth and Laurillard 2017). The developmental dyscalculia analysis (Supplementary Material 5) is an early step in this direction: the same model localises the deficit not as noisier perception but as a shorter exact range, a code whose resolution runs out earlier.

The approximate regime is the same landscape read where the wells have merged. Beyond n_{crit} clean-up no longer sharpens the posterior, and the same readout

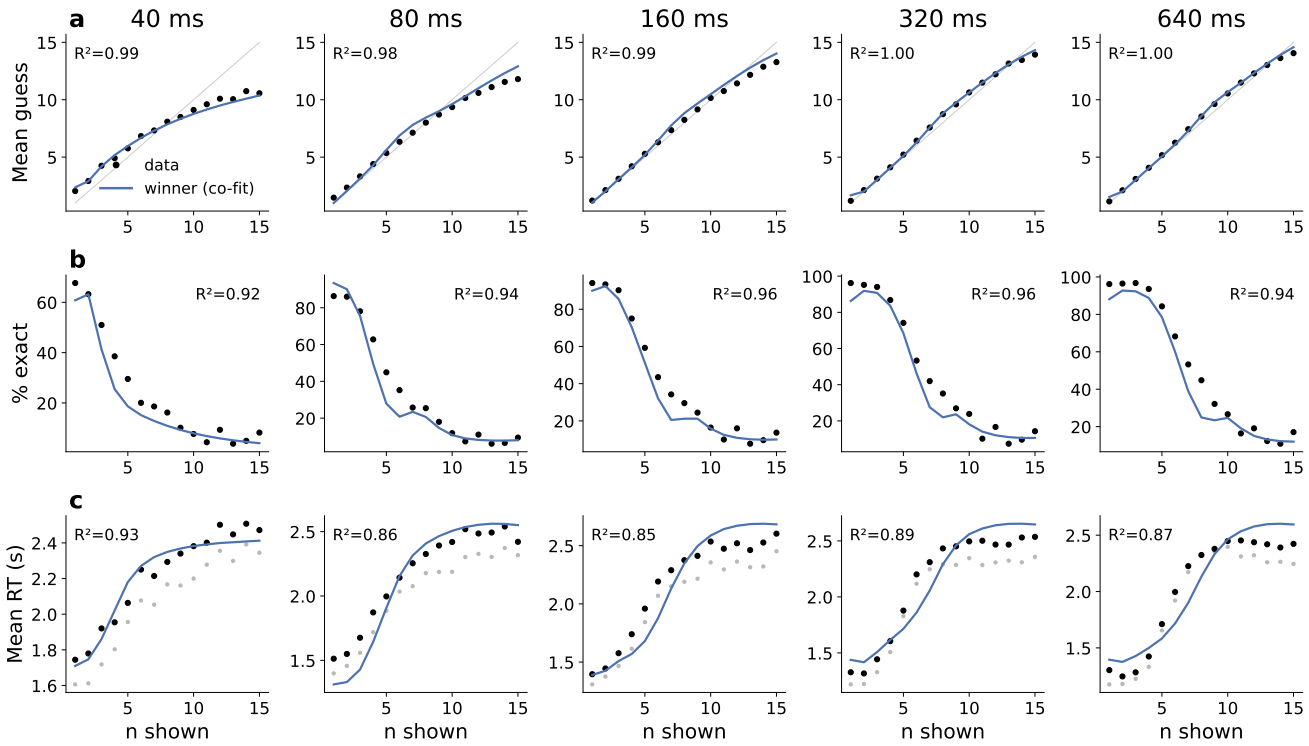


Figure 3. Joint fit of the model to the Cheyette and Piantadosi (2020) data. Columns show presentation times from 40 to 640 ms; rows show the mean estimate, exact accuracy $P(r = n)$, and geometric-mean response time. Dots are empirical data (faint dots in the RT column are the median, shown for reference); lines are the model, co-fitted across all three channels. Response time is a cross-prediction: the RT channel is fitted to the individual trial times, not to the curve shown. R^2 values are given per panel.

returns a coarse, ratio-dependent estimate, the behaviour attributed to the Approximate Number System. This needs no separate module: the model's response distributions match the data across the whole range, and their coefficient of variation settles to a roughly constant plateau across the estimation range, matching the Weber fraction measured in numerosity estimation (Testolin and McClelland 2021) (Supplementary Material 4). One mechanism therefore accounts for numerosity estimation in both regimes, exact where the wells resolve and Weber-law approximate where they merge.

What the model does not capture is chronometric. Within the exact range the convergence-time term tracks response time, but beyond n_{crit} it keeps rising while the empirical RT plateaus (Fig. 1c). The model runs clean-up to convergence at every numerosity; participants appear instead to stop deliberating once the code is no longer resolvable, reading the merged posterior directly and answering quickly. Some mechanism halts the clean-up around n_{crit} and flattens the response time, and what triggers it is unknown. Because stopping early and running to convergence leave a similarly broad readout on merged wells, this affects the latency but not the estimate, which is why the model captures the approximate estimates while over-predicting how long they take. Whether numerosity adaptation involves the same halting is a natural question for future work.

This account offers a different level of explanation from the leading unifying alternative. Cheyette et al. (2024) derive number psychophysics from a bandwidth limit on visuospatial memory; in our model that limit would live in the sensory channel $\sigma(n)$, upstream of the attractor, and the observation that n_{crit} grows with exposure is consistent with more available bandwidth yielding a finer code. The accounts are therefore complementary rather than rival, and both invoke a frequency prior of the Benford family. Where they differ is in what the number-specific representation contributes: our attractor is modality-general (any channel delivering a scalar estimate feeds the same wells) and, being a memory, is intrinsically trainable. It also differs from recurrent accounts that partition the numerosity range through fixed network dynamics (Sengupta et al. 2014): our boundary is not a fixed partition but one that moves with exposure and, we predict, with development, a discriminating test between the two.

Several limitations remain. The evidence channel is deliberately neutral: it summarises the visual front end as a noise profile $\sigma(n)$, so the model operates on an already-extracted estimate. This isolates the attractor dynamics but, by construction, sets aside the influence of continuous magnitudes (area, density, convex hull) on numerosity judgements; coupling the wells to an image-computable front end that is itself sensitive to number (Stoianov and Zorzi 2012; Testolin et al. 2020), and that has been used to simulate dyscalculic profiles (Dolfi et al. 2024), is the clearest way to close this gap. The model also under-predicts the sharpness of small-number accuracy at $n =$ 2–4: it captures the trend but not the full abruptness of the subitizing edge, which may still require a fast parallel individuation process for the very smallest sets (Feigenson et al. 2002). Finally, the model is specified at a computational and algorithmic level; relating the wells to the number-tuned populations of parietal cortex, whose tuning widths grow

with preferred numerosity (Piazza et al. 2004; Harvey et al. 2013), is a natural direction for testing it neurally.

Numerical cognition, in this framework, is pattern completion on a landscape of number memories. A single set of wells, read exactly where they are resolved and approximately where they merge, accounts for the fast exactness of small numbers, the graded approximation of larger ones, the response times of both, and the moving boundary between them, offering a compact and mechanistic set of hypotheses about how the mind represents number.

Acknowledgements

I am grateful to Brian Butterworth, Diana Laurillard and Marc Colomer for their insightful discussions and thoughtful feedback, which substantially improved this work.

Author Contributions

Eudald Correig-Fraga: Conceptualization, methodology, software, formal analysis, investigation, data curation, visualization, writing - original draft, and writing - review and editing.

Ethical Considerations

This article reports computational modelling and secondary analysis of previously published, publicly available datasets. No new human participants were recruited, and no new human data were collected for this study. Ethical approval for the original data collections was reported in the source publications.

Consent to Participate

Not applicable. This study did not recruit participants or collect new participant data.

Consent for Publication

Not applicable. This study does not contain identifiable individual participant information.

Funding

This work was funded by Innovamat Education SL.

Declaration of conflicting interests

The author declared no potential conflicts of interest.

Data Availability

Code is available at https://github.com/eudaldc/hopfield_numbers. The source datasets are available at https://osf.io/svcy5/ and https://osf.io/gjpb3/.

Supplemental material

Supplementary material is available with this article.

References

Anderson JR and Schooler LJ (1991) Reflections of the Environment in Memory. *Psychological Science* 2(6): 396–408. DOI:10.1111/j.1467-9280.1991.tb00174.x.

Butterworth B and Laurillard D (2017) Investigating dyscalculia: A science of learning perspective. In: *From the laboratory to the classroom: Translating science of learning for teachers*. New York, NY, US: Routledge/Taylor & Francis Group. ISBN 978-1-138-64964-4 978-1-138-64963-7 978-1-315-62573-7, pp. 172–190.

Cheyette SJ and Piantadosi ST (2020) A unified account of numerosity perception. *Nature Human Behaviour* 4(12): 1265–1272. DOI:10.1038/s41562-020-00946-0.

Cheyette SJ, Wu S and Piantadosi ST (2024) Limited information-processing capacity in vision explains number psychophysics. *Psychological Review* 131(4): 891–904. DOI:10.1037/rev0000478.

Decarli G, Paris E, Tencati C, Nardelli C, Vescovi M, Surian L and Piazza M (2020) Impaired large numerosity estimation and intact subitizing in developmental dyscalculia. *PLOS ONE* 15(12): e0244578. DOI:10.1371/journal.pone.0244578.

Dehaene S (2011) *The number sense: How the mind creates mathematics, Rev. and updated ed*. New York, NY, US: Oxford University Press. ISBN 978-0-19-975387-1. Pages: xxii, 316.

Dehaene S and Mehler J (1992) Cross-linguistic regularities in the frequency of number words. *Cognition* 43(1): 1–29. DOI:10.1016/0010-0277(92)90030-L.

Dolfi S, Decarli G, Lunardon M, De Filippo De Grazia M, Gerola S, Lanfranchi S, Cossu G, Sella F, Testolin A and Zorzi M (2024) Weaker number sense accounts for impaired numerosity perception in dyscalculia: Behavioral and computational evidence. *Developmental Science* 27(6): e13538. DOI:10.1111/desc.13538. eprint: https://onlinelibrary.wiley.com/doi/pdf/10.1111/desc.13538.

Feigenson L, Carey S and Hauser M (2002) The representations underlying infants' choice of more: object files versus analog magnitudes. *Psychological Science* 13(2): 150–156. DOI:10.1111/1467-9280.00427.

Feigenson L, Dehaene S and Spelke E (2004) Core systems of number. *Trends in Cognitive Sciences* 8(7): 307–314. DOI:10.1016/j.tics.2004.05.002.

Halberda J, Mazzocco MM and Feigenson L (2008) Individual differences in non-verbal number acuity correlate with maths achievement. *Nature* 455(7213): 665–668. DOI:10.1038/nature07246.

Harvey BM, Klein BP, Petridou N and Dumoulin SO (2013) Topographic Representation of Numerosity in the Human Parietal Cortex. *Science* 341(6150): 1123–1126. DOI:10.1126/science.1239052.

Hopfield JJ (1982) Neural networks and physical systems with emergent collective computational abilities. *Proceedings of the National Academy of Sciences* 79(8): 2554–2558. DOI:10.1073/pnas.79.8.2554.

Kaufman EL, Lord MW, Reese TW and Volkmann J (1949) The Discrimination of Visual Number. *The American Journal of Psychology* 62(4): 498–525. DOI:10.2307/1418556.

Macmillan NA and Creelman CD (2004) *Detection Theory: A User's Guide*. 2 edition. New York: Psychology Press. ISBN 978-1-4106-1114-7. DOI:10.4324/9781410611147.

Malo JY, Cicchini GM, Morrone MC and Chiofalo ML (2023) Quantum spin models for numerosity perception. *PLOS ONE* 18(4): e0284610. DOI:10.1371/journal.pone.0284610. ArXiv:2212.03344 [q-bio.NC].

Mandler G and Shebo BJ (1982) Subitizing: An analysis of its component processes. *Journal of Experimental Psychology: General* 111(1): 1–22. DOI:10.1037/0096-3445.111.1.1.

Nieder A (2016) The neuronal code for number. *Nature Reviews Neuroscience* 17(6): 366–382. DOI:10.1038/nrn.2016.40.

Piazza M, Fumarola A, Chinello A and Melcher D (2011) Subitizing reflects visuo-spatial object individuation capacity. *Cognition* 121(1): 147–153. DOI:10.1016/j.cognition.2011.05.007.

Piazza M, Izard V, Pinel P, Le Bihan D and Dehaene S (2004) Tuning curves for approximate numerosity in the human intraparietal sulcus. *Neuron* 44: 547–555.

Pica P, Lemer C, Izard V and Dehaene S (2004) Exact and approximate arithmetic in an Amazonian indigene group. *Science* 306(5695): 499–503. DOI:10.1126/science.1102085.

Ramsauer H, Schäfl B, Lehner J, Seidl P, Widrich M, Adler T, Gruber L, Holzleitner M, Pavlović M, Sandve GK, Greiff V, Kreil D, Kopp M, Klambauer G, Brandstetter J and Hochreiter S (2021) Hopfield Networks is All You Need. In: *International Conference on Learning Representations*.

Sengupta R, Surampudi BR and Melcher D (2014) A visual sense of number emerges from the dynamics of a recurrent on-center off-surround neural network. *Brain Research* 1582: 114–124.

Stoianov I and Zorzi M (2012) Emergence of a 'visual number sense' in hierarchical generative models. *Nature Neuroscience* 15(2): 194–196. DOI:10.1038/nn.2996.

Storn R and Price K (1997) Differential Evolution - A Simple and Efficient Heuristic for Global Optimization over Continuous Spaces. *Journal of Global Optimization* 11: 341–359. DOI:10.1023/A:1008202821328.

Testolin A, Dolfi S, Rochus M and Zorzi M (2020) Visual sense of number vs. sense of magnitude in humans and machines. *Scientific Reports* 10(1): 10045. DOI:10.1038/s41598-020-66838-5.

Testolin A and McClelland JL (2021) Do estimates of numerosity really adhere to Weber's law? A reexamination of two case studies. *Psychonomic Bulletin & Review* 28(1): 158–168. DOI:10.3758/s13423-020-01801-z.

Verguts T and Fias W (2004) Representation of Number in Animals and Humans: A Neural Model. *Journal of Cognitive Neuroscience* 16(9): 1493–1504.